

▼ Daniel Vargas Diaz

Homework #4 CS 5844: Human-Robot Interaction

▼ Predict and Blend

▼ 1.1

Consider a shared autonomy algorithm that blends the robot action a_R and human input a_H according to:

$$a = (1 - \alpha) \cdot a_H + \alpha \cdot a_R \quad (1)$$

where $\alpha \in (0, 1)$. The human takes actions that move directly towards the goal and the robot takes actions in the opposite direction with magnitude Δ :

$$a_H = \theta^* - s, \quad a_R \propto -(\theta^* - s), \quad \|a_R\| = \Delta \quad (2)$$

Find the distance between the state s and the human's goal θ^* when the system converges. Your answer should be a function of α and Δ .

We know from the question that

$$\Delta = \|a_R\|$$

We could say that the distance between the goal θ^* and the state s multiply by a constant K is equal to the magnitude of the robot action

$$\Delta = K * \|\theta^* - s\|$$

$$K = \frac{\Delta}{\|\theta^* - s\|}$$

We also know that the system converges when the action taken by the system is equal to 0

$$0 = (1 - \alpha) * a_H + \alpha * a_R$$

$$0 = (1 - \alpha) * \theta^* - s - \alpha * K * \theta^* - s$$

$$0 = (1 - \alpha) * (\theta^* - s) - \alpha * K * (\theta^* - s)$$

$$0 = (\theta^* - s) * (1 - \alpha - \alpha * K)$$

From this we have two equations. If we take the first term, equal to 0, we said that θ^* equal to s , which does not make sense, this means, that the other term is the one that should be equal to zero

$$0 = (1 - \alpha - \alpha * K)$$

$$0 = (1 - \alpha - \frac{\alpha * \Delta}{\|\theta^* - s\|})$$

$$\|\theta^* - s\| = \frac{\alpha * \Delta}{(1 - \alpha)}$$

▼ 1.2

Download the provided code predict-blend.py. Here a simulated human may change their goal halfway through the task, and we need to help them reach their goal despite this change. Develop your own approach for predicting the human's goal and assisting the human's actions. Describe your approach using pseudocode. Across 1000 runs, show that your approach results in an average error of < 0.4 units

```

1 """
2 Starter code for HW4 Problem 1.2
3
4 """
5
6 import numpy as np
7
8
9 def predict_goal(s0, st, THETA, prior, beta=1):
10     P = [0.] * len(THETA)
11     for idx in range(len(THETA)):
12         P[idx] = prior[idx]
13     dist_so_far = np.linalg.norm(st - s0)
14     for idx, theta in enumerate(THETA):
15         dist_to_goal = np.linalg.norm(theta - st)
16         min_dist = np.linalg.norm(theta - s0)
17         num = np.exp(beta * min_dist)
18         den = np.exp(beta * (dist_so_far + dist_to_goal))
19         P[idx] *= num / den
20     P = np.asarray(P)
21     return P / np.sum(P)
22
23 def get_assist(s, THETA, P):
24     aR = np.array([0., 0.])
25     for idx, theta in enumerate(THETA):
26         atheta = theta - s
27         if np.linalg.norm(atheta) > 1.0:
28             atheta /= np.linalg.norm(atheta)
29         aR += P[idx] * atheta
30     return aR
31
32 # simulated human
33 # this human may change their goal midway through the task
34 def simulated_human(s, t, human_goal, THETA):
35     # at timestep t = 10, the human might switch goals
36     if t == 10:
37         goal_idx = np.random.choice(2)
38         human_goal = THETA[goal_idx]
39     # human takes noisy action towards chosen goal
40     ah = human_goal - s + np.random.normal(0, 0.3, 2)
41     if np.linalg.norm(ah) > 1.0:
42         ah /= np.linalg.norm(ah)
43     return ah, human_goal
44
45 # main loop; each run has new goals and a new start position
46 def main():
47
48     # each run we have two randomly placed goals
49     theta1 = 10.0 * np.random.rand(2)
50     theta2 = 10.0 * np.random.rand(2)
51     THETA = [theta1, theta2]
52     human_goal = THETA[np.random.choice(2)]
53
54     # robot starts in a random state
55     s = 10.0 * np.random.rand(2)
56
57     # robot starts with uniform belief over the goals
58     b = np.array([0.5, 0.5])
59
60     # initialize state
61     s0 = np.copy(s)
62
63     # for T=20 timesteps
64     #print("Initial Goal: ", human_goal)
65     #print("Initial State:", s0)
66     for t in range(20):
67
68         # get human action

```

```

69     # human_goal is the human's actual goal
70     # we only use this to check our approach's accuracy;
71     # the robot does not have access to "human_goal"
72     ah, human_goal = simulated_human(s, t, human_goal, THETA)
73     #print("Human action: ",ah)
74
75     # implement your algorithm here for helping the human
76     # your approach should predict the human's goal
77     # then take actions to assist
78     # remember: you know THETA, but not human_goal
79     #p_theta = predict_goal(s0, s, THETA, b)
80     #p_theta_all.append(p_theta)
81     if(t==10):
82         #print("can change now")
83         s0 = np.copy(s)
84     if(t<11):
85         ar = get_assist(s, THETA, b)
86     elif(t>=11):
87         p_theta = predict_goal(s0, s, THETA, b)
88         ar = get_assist(s, THETA, p_theta)
89         if(p_theta[0]>0.6 or p_theta[1]>0.6):
90             b = p_theta
91
92
93     #print("State: ", s)
94
95     # take blended action
96     alpha = 0.8 # do not change this value
97     a = (1 - alpha) * ah + alpha * ar
98
99     # transition to the next state
100    s1 = s + a
101    s = np.copy(s1)
102    #print("Final Goal: ", human_goal)
103    #print("Final state: ", s)
104    #print("Error: ", s - human_goal)
105    return np.linalg.norm(s - human_goal)
106
107 # run the main loop and find average error
108 total_error = 0.
109 for _ in range(1000):
110     error = main()
111     total_error += error
112 print(total_error / 1000.)
113
0.3144225444575468

```


▼ Partially Observable Markov Decision Process

Consider a 5×5 gridworld where the robot starts at $(0,0)$ and the human is standing at $(4,4)$. The robot can move up, down, left, right, or *stay in place*, and the dynamics are deterministic. The robot is trying to quickly rescue an item from the room, but it does not know where the item is. With 70% probability the item is at position $(3,1)$, and with 30% probability the item is at $(1,3)$. The robot can ask the human and remove any uncertainty by moving to position $(4,4)$. The reward at every state besides the item's position is 0.

▼ 2.1

Let the reward for reaching the item be +10.0. Is it optimal for the robot to ask to the human if $\gamma = 0.5$? Draw the robot's trajectory starting from $(0,0)$.

In this case, the robot consider that is not optimal to ask the human. This is because the discount factor is small enough to make the robot only search for the direct reward, which in this case is the most likely reward, the one with 70% of chances.



```

1 """
2 Goal POMDP (example)
3
4 """
5
6 import numpy as np
7
8
9 class POMDP:
10
11     # initialization
12     def __init__(self):
13
14         # state space
15         # here states include x-y position
16         # and the belief that the goal is at theta_1
17         # b = 1 means we are confident the goal is at theta_1
18         # b = 0 means we are confident the goal is at theta_2
19         self.states = []
20         for s_x in range(0, 5):
21             for s_y in range(0, 5):
22                 for b in np.linspace(0, 1, 4):
23                     b = round(b, 1)
24                     self.states.append((s_x, s_y, b))
25
26         # action space
27         self.actions = ((0,1),(0,-1),(1,0),(-1,0),(0,0))
28         # discount factor
29         self.gamma = 0.5
30
31     # deterministic dynamics
32     def f(self, s, a):
33
34         # take commanded action
35         next_states = {}

```

```


36     s1 = self.take_action(s, a)
37
38     # get human input
39     # this occurs at a specific state
40     # only at this state do we get insight from the
41     # human about which goal to go for :)
42     if s[0] == 4 and s[1] == 4:
43         s1_goal1 = (s1[0], s1[1], 1.0)
44         s1_goal2 = (s1[0], s1[1], 0.0)
45         # 70% of the time is going to be theta_1
46         # 30% of the time is going to be theta_2
47         next_states[s1_goal1] = 0.7
48         next_states[s1_goal2] = 0.3
49     else:
50         next_states[s1] = 1.0
51
52     return next_states
53
54 # helper function for dynamics
55 def take_action(self, s, a):
56     # s1 = s + a, belief stays constant
57     s1 = (s[0] + a[0], s[1] + a[1], s[2])
58     # keeps robot inside state space
59     if s1 in self.states:
60         return s1
61     else:
62         return s
63
64 # reward for goal 1
65 # this goal is at position (3, 1)
66 def reward1(self, s):
67     if s[0] == 3 and s[1] == 1:
68         return +10.0
69     else:
70         return 0.0
71
72 # reward for goal 2
73 # this goal is at position (1, 3)
74 def reward2(self, s):
75     if s[0] == 1 and s[1] == 3:
76         return +10.0
77     else:
78         return 0.0
79
80
81 # Q-function
82 def Qfunction(mdp, s, a, V):
83     next_states = mdp.f(s, a)
84     expected_V1 = sum(next_states[s1] * V[s1] for s1 in next_states)
85     # get expected reward using belief
86     expected_reward = s[2] * mdp.reward1(s) + (1.0 - s[2]) * mdp.reward2(s)
87     return expected_reward + mdp.gamma * expected_V1
88
89 # get optimal policy from value function
90 def policy(mdp, V):
91     pi = {}
92     for s in mdp.states:
93         # take action that maximizes Q-function
94         max_Q = -np.inf
95         for action in mdp.actions:
96             Q = Qfunction(mdp, s, action, V)
97             if Q > max_Q:
98                 max_Q = Q
99             pi[s] = action
100     return pi
101
102 # value iteration algorithm to get value function
103 def value_iteration(mdp, epsilon=0.001):

```


▼ 2.2

Let the reward for reaching the item be +1.0. Is it optimal for the robot to ask to the human if $\gamma = 0.99$? Draw the robot's trajectory starting from (0, 0).

In this case it is optimal for the robot to ask the human, because the discount factor is big enough to make the robot focus on the long term reward. This results in the robot asking the human and the going to one of the goals.



```

1 """
2 Goal POMDP (example)
3
4 """
5
6 import numpy as np
7
8
9 class POMDP:
10
11     # initialization
12     def __init__(self):
13
14         # state space
15         # here states include x-y position
16         # and the belief that the human wants theta_1
17         # b = 1 means we are confident human wants theta_1
18         # b = 0 means we are confident human wants theta_2
19         self.states = []
20         for s_x in range(0, 5):
21             for s_y in range(0, 5):
22                 for b in np.linspace(0, 1, 4):
23                     b = round(b, 1)
24                     self.states.append((s_x, s_y, b))
25
26         # action space
27         self.actions = ((0,1),(0,-1),(1,0),(-1,0),(0,0))
28         # discount factor
29         self.gamma = 0.99

```

```

30
31 # deterministic dynamics
32 def f(self, s, a):
33
34     # take commanded action
35     next_states = {}
36     s1 = self.take_action(s, a)
37
38     # get human input
39     # this occurs at a specific state
40     # only at this state do we get insight from the
41     # human about which goal to go for :)
42     if s[0] == 4 and s[1] == 4:
43         s1_goal1 = (s1[0], s1[1], 1.0)
44         s1_goal2 = (s1[0], s1[1], 0.0)
45         # half the time the human tells us theta_1
46         # the other half the human tells us theta_2
47         next_states[s1_goal1] = 0.3
48         next_states[s1_goal2] = 0.7
49     else:
50         next_states[s1] = 1.0
51
52     return next_states
53
54 # helper function for dynamics
55 def take_action(self, s, a):
56     # s1 = s + a, belief stays constant
57     s1 = (s[0] + a[0], s[1] + a[1], s[2])
58     # keeps robot inside state space
59     if s1 in self.states:
60         return s1
61     else:
62         return s
63
64 # reward for goal 1
65 # this goal is at position (3, 1)
66 def reward1(self, s):
67     if s[0] == 3 and s[1] == 1:
68         return +1.0
69     else:
70         return 0.0
71
72 # reward for goal 2
73 # this goal is at position (1, 3)
74 def reward2(self, s):
75     if s[0] == 1 and s[1] == 3:
76         return +1.0
77     else:
78         return 0.0
79
80
81 # Q-function
82 def Qfunction(mdp, s, a, V):
83     next_states = mdp.f(s, a)
84     expected_V1 = sum(next_states[s1] * V[s1] for s1 in next_states)
85     # get expected reward using belief
86     expected_reward = s[2] * mdp.reward1(s) + (1.0 - s[2]) * mdp.reward2(s)
87     return expected_reward + mdp.gamma * expected_V1
88
89 # get optimal policy from value function
90 def policy(mdp, V):
91     pi = {}
92     for s in mdp.states:
93         # take action that maximizes Q-function
94         max_Q = -np.inf
95         for action in mdp.actions:
96             Q = Qfunction(mdp, s, action, V)
97             if Q > max_Q:

```



```

98         max_Q = Q
99         pi[s] = action
100     return pi
101
102 # value iteration algorithm to get value function
103 def value_iteration(mdp, epsilon=0.001):
104     V1 = {s: 0 for s in mdp.states}
105     while True:
106         V = V1.copy()
107         delta = 0
108         for s in mdp.states:
109             # use value function V to get estimate V'
110             action_values = []
111             for action in mdp.actions:
112                 next_states = mdp.f(s, action)
113                 action_value = sum(next_states[s1] * V[s1] for s1 in next_states)
114                 action_values.append(action_value)
115             # get expected reward using belief
116             expected_reward = s[2] * mdp.reward1(s) + (1.0 - s[2]) * mdp.reward2(s)
117             V1[s] = expected_reward + mdp.gamma * max(action_values)
118             # find largest change in value
119             delta = max(delta, abs(V1[s] - V[s]))
120         # stop if largest change is small enough
121         if delta <= epsilon * (1 - mdp.gamma) / mdp.gamma:
122             return V
123
124 def main():
125     # setup the pomdp
126     # this functions the same as an mdp
127     # but now with belief as part of state
128     robot = POMDP()
129
130     # use value iteration to get V
131     V = value_iteration(robot)
132
133     # get optimal policy from V
134     pi = policy(robot, V)
135
136     # print policy
137     print(pi)
138
139     # rollout policy from chosen start state
140     s = (0, 0, 0.3)
141     for t in range(20):
142         a = pi[s]
143         print(s, a)
144         next_states = robot.f(s,a)
145         # sorry this part is confusing
146         # need to pick the next state
147         s = list(next_states.keys())[0]
148
149
150 main()
151
152 {(0, 0, 0.0): (0, 1), (0, 0, 0.3): (0, 1), (0, 0, 0.7): (0, 1), (0, 0, 1.0): (0, 1), (0, 1, 0.0): (0, 1), (0, 1, 0.3):
(0, 1), (0, 1, 0.3): (0, 1), (0, 2, 0.3): (0, 1), (0, 3, 0.3): (1, 0), (1, 3, 0.3): (0, 1), (1, 4, 0.3): (1, 0), (2, 4, 0.3): (1, 0), (3, 4, 0.3): (1, 0), (4, 4, 0.3): (0, -1), (4, 3, 1.0): (0, -1), (4, 2, 1.0): (0, -1), (4, 1, 1.0): (-1, 0), (3, 1, 1.0): (0, 0), (3, 1, 1.0): (0, 0)}

```

```
(3, 1, 1.0) (0, 0)
(3, 1, 1.0) (0, 0)
(3, 1, 1.0) (0, 0)
(3, 1, 1.0) (0, 0)
(3, 1, 1.0) (0, 0)
(3, 1, 1.0) (0, 0)
```

▼ 3. Latent Actions

▼ 3.1

Imagine that you have a matrix of 7-dimensional robot actions. You want to control the robot using 2-DoF. You propose finding these embeddings through singular value decomposition. Write detailed pseudocode for controlling the robot using the first two eigenvectors.

Assuming I have 2 DoF, this means I have only 2 axis, in this case I am going to say X,Y.

X+ -> Output_1: When we move the X vector to positive direction we control the Output 1

X- -> Output_2: When we move the X vector to negative direction we control the Output 2

Y+ -> Output_3: When we move the Y vector to positive direction we control the Output 3

Y- -> Output_4: When we move the X vector to negative direction we control the Output 4

Sin(citha) -> Output_5: The Sin of the angle citha, that is the angle between X and Y, control the Output 5

Cos(citha) -> Output_6: The Cos of the angle citha, that is the angle between X and Y, control the Output 6

Tan(citha) -> Output_7: The Tan of the angle citha, that is the angle between X and Y, control the Output 7

3.2 Now consider the opposite direction; you have a 7-dimensional input device and are controlling a 2-DoF robot. Describe your own algorithm for mapping high-dimensional inputs to low-dimensional outputs. Explain why your approach would be intuitive for the user.

Assuming that I have 7 input device, like a control with 4 arrows and 3 buttons, to control 2 DoF (X,Y), I would do this:

Map the Input 1 to control X positive direction with the current citha angle that is the angle between X and Y

Map the Input 2 to control X negative direction with the current citha angle that is the angle between X and Y

Map the Input 3 to control Y positive with the current citha angle that is the angle between X and Y

Map the Input 4 to control transforms Y negative direction with the current citha angle that is the angle between X and Y

Map the Input 5 to transform the Input 1, 2, 3 and 4 to control the citha angle that is the angle between X and Y from 0 to 180

Map the Input 6 to transform the Input 1, 2, 3 and 4 to control the citha angle that is the angle between X and Y from 180 to 360

Map the Input 7 to activate the angle movement (5-6 become usable) or only use the 1D movement of Inputs 1,2,3 and 4

In this case would be intuitive because with the arrows they can move normally the robot, and with the 4 and 5 buttons they can activate the function to convert the arrows into mappings of 1D to Angle vectors, and with the other button they can activate or deactivate this mapping angle functions. With this approach the arrows will be only managing a specific task depending of the combination of 5, 6 and 7

✓ 8 s se ejecutó 17:04

● ×